Object Detection: VOC2007

Overview:

The goal of this challenge is to recognize objects from a number of visual object classes in realistic scenes (i.e. not pre-segmented objects). It is fundamentally a supervised learning problem in that a training set of labeled images is provided. The twenty object classes that have been selected are:

- *Person:* person
- *Animal:* bird, cat, cow, dog, horse, sheep
- *Vehicle:* aeroplane, bicycle, boat, bus, car, motorbike, train
- *Indoor:* bottle, chair, dining table, potted plant, sofa, tv/monitor

In this Hackathon, your task is to build a CNN Object detector capable of classifying and detecting images

Data:

Download dataset from here http://pjreddie.com/media/files/VOCtrainval_06-Nov-2007.tar

Use the following resource to get started https://github.com/Anil-matcha/YOLOv2/blob/master/YOLOv2.ipynb

Public and Private LeaderBoard

- Your initial responses will be checked and scored on the Public data.
- The final rankings would be based on your private score which will be published once the competition is over.

Evaluation Criteria

Your model performance will be evaluated on the basis of the **MAP Score**.

Rubric

Component	Weightage
Data Cleaning and Data Visualization	25%
Model Building and Evaluation	60%
Pipeline and Deployment (Dashboard/Webapp)	15%